Earnings Surprise, Firm Valuation, and Market Efficiency: Evidence from China

Corporate Finance Team Project

Team Members: Yingqi Tan, Zhesheng Xu, Hao Yang

FIN803 Corporate Finance

Motivation

- Earnings-related disclosures update expectations about future cash flows and should affect valuation.
- The current goal is not to force a strong pooled headline, but to build a cleaner diagnostic baseline.
- Tushare Pro lets us test how expectation measurement and match quality affect short-run inference.

Research Question

Core Question

Do China A-share earnings-related events generate short-run abnormal return once surprise is measured relative to stricter sell-side expectations, and how much do signal scaling and match quality change that inference?

- Old version: guidance-only proxy design.
- Current version: Tushare-first diagnostic baseline with explicit event types and explicit match tiers.
- Final framing now emphasizes identification quality, not a forced positive finding.

New Headline Baseline

- Headline sample: `preannouncement_only`
- Match tier: `strict_same_quarter`
- Windows: `CAR3`, `CAR5`, `CAR10`, `CAR20`
- Compare `raw`, `pct`, and `std` in parallel

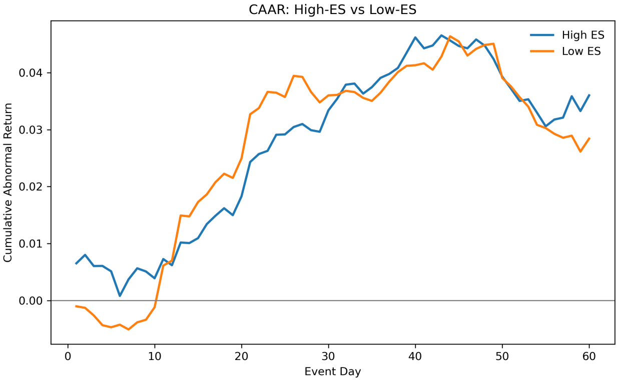

Diagnostic example: cumulative returns for high- vs low-surprise groups.

Why We Changed the Headline

- A pooled all-event standardized specification hides important heterogeneity.
- Preannouncements are cleaner than revisions, express releases, and formal releases for a default headline sample.
- `main_surprise_std` should be a robustness lens, not the only lens.
- Match-quality tiers should be explicit rather than silently blended.

Expectation-Matching Ladder

- `strict_same_quarter`
- `same_fiscal_year_nearest_valid`
- `latest_valid_pre_event`
- `multi_report_median`

Interpretation Rule

The headline baseline uses only `strict_same_quarter`; relaxed tiers are reported separately to show how benchmark construction changes inference.

Focused Diagnostic Package

- `preannouncement_only` vs `all_event_types`
- `raw` vs `pct` vs `std`
- `CAR3`, `CAR5`, `CAR10`, `CAR20`
- strict vs relaxed matching tiers
- tier_1 baseline vs tier_2 / tier_3 augmentation
- minimum analyst coverage thresholds

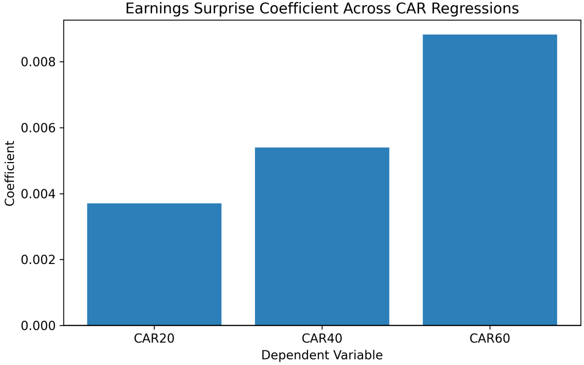

Regression-based view of how the measured surprise signal moves with specification choices.

Free-data Augmentation as Diagnostic Extension

- We keep `tier_1_tushare_report_rc` as the default headline baseline.
- We add `tier_2_eastmoney_profit_forecast` and `tier_3_eastmoney_research_report_text` only as labeled augmentation tiers.
- The goal is to study coverage-versus-quality tradeoffs, not to claim that free public proxies solve the expectation-quality problem.
- The main narrow comparison remains preannouncement-only, raw surprise, and CAR10-focused diagnostics.

How to Read Augmentation Results

Interpretation rule

If a free-data tier adds coverage but weakens identification materially, it should remain a supplementary diagnostic. Only if coverage expands without obvious deterioration should it be shown as an augmentation route — not as a replacement baseline.

What the Repository Claims Now

Recommended narrative

This repository is best viewed as a Tushare-first diagnostic baseline with a controlled free-data augmentation layer. The augmentation layer is used to test coverage-quality tradeoffs, not to replace the core baseline or to overstate headline evidence.

Updated Recommendation

Recommendation B+

Keep Tushare report_rc as the default expectation benchmark, keep preannouncement-only strict matching as the headline sample, and use Eastmoney tiers only as explicit diagnostic extensions when evaluating whether additional free coverage is worth the loss in expectation quality.

Next Step

- Keep the Tushare event architecture and focused diagnostic outputs.
- Use preannouncement-only strict matching with tier_1 as the baseline comparison set.
- Evaluate free-data augmentation only through separate labeled tier comparisons.
- Add stronger external analyst-expectation data only if the goal is a stronger final headline.

Conclusion

- We kept the Tushare-first architecture and narrowed the headline sample.
- We now compare raw, pct, and std in parallel instead of relying on one standardized headline.
- We now expose strict and relaxed match tiers explicitly.
- We add free-data augmentation as a diagnostic extension rather than a replacement.
- The repository's main contribution is still a cleaner and more auditable diagnostic baseline.