Speaker Script
Earnings Surprise, Firm Valuation, and Market Efficiency: Evidence from China

FIN803 Corporate Finance Team

Slide 1: Title

Good afternoon everyone. We are presenting our FIN803 Corporate Finance project on earnings surprise, valuation adjustment, and market efficiency in China A-share stocks.

Slide 2: Motivation

Our motivation is to move away from a weak pooled headline and toward a cleaner Tushare-first diagnostic baseline. The key question is not whether one broad specification looks significant, but how short-run inference changes when expectation measurement and match quality become more disciplined.

Slide 3: Research Question

Our research question is now framed more narrowly. We ask whether China A-share earnings-related events generate short-run abnormal return when surprise is measured against stricter sell-side expectations, and how much that answer changes across signal scaling choices and explicit match tiers.

Slide 4: New Headline Baseline

The default headline sample is now preannouncement-only with strict same-quarter matching. We compare CAR3, CAR5, CAR10, and CAR20, and we compare raw, percent, and standardized surprise in parallel. The figure on this slide gives a visual diagnostic example by contrasting the cumulative return path of high-surprise and low-surprise groups.

Slide 5: Why We Changed the Headline

We made this change because the pooled all-event standardized design hides important heterogeneity. Preannouncements are cleaner as a default headline sample, and standardization should be treated as a robustness lens rather than the only lens.

Slide 6: Expectation-Matching Ladder

We also rebuilt the expectation-matching logic into an explicit ladder. The tiers are strict same quarter, same fiscal year nearest valid, latest valid pre-event, and multi-report median. The important point is that we no longer silently blend these tiers. The strict tier is the headline baseline, and the relaxed tiers are reported separately as diagnostics.

Slide 7: Focused Diagnostic Package

The repository now produces a side-by-side diagnostic package covering preannouncement-only versus all event types, raw versus percent versus standardized surprise, CAR3 through CAR20, strict versus relaxed matching, tier augmentation, and minimum analyst coverage thresholds. The regression figure on this slide adds a second lens for how the measured signal moves with specification choices.

Slide 8: Free-data Augmentation as Diagnostic Extension

A key extension is a controlled free-data augmentation layer. We keep Tushare report_rc as the default expectation baseline, and we add Eastmoney profit-forecast and research-report text only as labeled augmentation tiers. The purpose is to study coverage-versus-quality tradeoffs, not to claim that free public proxies solve the expectation-quality problem.

Slide 9: How to Read Augmentation Results

The interpretation rule is simple. If a free-data tier adds coverage but materially weakens identification, it should remain supplementary. Only if coverage expands without obvious deterioration should it be discussed as an augmentation route, and even then it should not replace the baseline by default.

Slide 10: What the Repository Claims Now

So the repository should now be read as a Tushare-first diagnostic baseline with a controlled augmentation layer. Its value is in showing how expectation measurement, source quality, and matching discipline shape the empirical conclusion, not in overstating headline evidence.

Slide 11: Updated Recommendation

Our updated recommendation is therefore a cautious B plus. Keep Tushare report_rc as the default expectation benchmark, keep preannouncement-only strict matching as the headline sample, and use Eastmoney tiers only as explicit diagnostic extensions when asking whether extra free coverage is worth the loss in expectation quality.

Slide 12: Next Step

The next step is not to redesign the framework again. The next step is to preserve the Tushare event architecture and focused diagnostics, evaluate free-data augmentation only through separate labeled tier comparisons, and add stronger external analyst-expectation data only if the goal is a stronger final headline.

Slide 13: Conclusion

Our final conclusion is narrower and more defensible. We kept the Tushare-first architecture, narrowed the headline sample, exposed match quality explicitly, and added free-data augmentation only as a diagnostic extension. The repository's main contribution remains a cleaner and more auditable diagnostic baseline.