Speaker Script
Earnings Surprise, Firm Valuation, and Market Efficiency: Evidence from China

FIN803 Corporate Finance Team

Slide 1: Title

Good afternoon everyone. We are presenting our FIN803 Corporate Finance project on earnings surprise, valuation adjustment, and market efficiency in China A-share stocks.

Slide 2: Motivation

Our motivation is to move away from a weak internal proxy design and toward a more defensible Tushare-first diagnostic baseline. The key question is not just whether some pooled specification is significant, but how inference changes when expectation measurement becomes cleaner and match quality is made explicit.

Slide 3: Research Question

Our research question is now framed more narrowly. We ask whether earnings-related events generate short-run abnormal return relative to stricter pre-event sell-side expectations, and how much that answer changes across raw, percent, and standardized surprise measures and across explicit matching tiers.

Slide 4: New Headline Baseline

The default headline sample is no longer the pooled all-event standardized specification. The default baseline is now preannouncement-only with strict same-quarter matching. We compare CAR3, CAR5, CAR10, and CAR20, and we compare raw, percent, and standardized surprise in parallel. The figure on this slide gives an intuitive diagnostic example by contrasting the cumulative return path of high-surprise and low-surprise groups.

Slide 5: Why We Narrowed the Headline

We made this change because the broader pooled design mixes economically different event families and can hide how much the conclusion depends on benchmark construction. Preannouncements are cleaner as a default headline sample, and standardization should be a robustness check rather than the only lens.

Slide 6: Matching Ladder

We also rebuilt the expectation-matching logic into an explicit ladder. The tiers are strict same quarter, same fiscal year nearest valid, latest valid pre-event, and multi-report median. The important point is that we no longer silently mix them. The strict tier is the headline baseline, and the relaxed tiers are supplementary diagnostics.

Slide 7: Focused Diagnostic Package

The repository now produces side-by-side tables for preannouncement-only versus all events, raw versus percent versus standardized surprise, CAR3 through CAR20, strict versus relaxed matching, and minimum analyst coverage thresholds. The figure on this slide adds a regression-based view of how measured earnings surprise changes across specifications.

Slide 8: Tushare-first Framework

The framework itself is still Tushare-first. We use report_rc for sell-side expectations, forecast and forecast_vip for management preannouncements, express and express_vip for earnings express releases, fina_indicator for formal realized data, and daily_basic for market controls and liquidity filters.

Slide 9: Why the Old Proxy Was Weak

The older guidance-only design inferred expectations from prior guidance, firm history, and industry heuristics. That is weaker than a direct sell-side benchmark and much harder to interpret cleanly. So the old path remains only as a legacy comparison, not as the default headline design.

Slide 10: How to Read the Current Repo

The current repository should be read as a diagnostic baseline, not as a strong positive result. Its main value is showing how event definition, surprise scaling, matching quality, and analyst coverage jointly shape the empirical conclusion.

Slide 11: What Still Survives

What still survives is narrow. The cleanest default sample is preannouncement-only. The cleanest default match tier is strict same-quarter. Raw surprise may still be the best main display specification if it remains the strongest surviving Tushare-only signal, while percent and standardized surprise provide parallel checks. The accompanying distribution figure is a reminder that the signal remains informative but should still be read cautiously.

Slide 12: Why We Stay Cautious

That surviving evidence should still be described as diagnostic rather than decisive. The current Tushare-only sample can be informative, but it does not justify a strong headline claim on its own. We should not overstate what the data can currently support.

Slide 13: Recommendation

So our recommendation remains cautious. Keep Tushare as the baseline framework and use the repository to diagnose how inference changes with signal scaling and match quality. If the goal is a stronger final empirical headline, stronger external analyst-expectation data are still the credible next step.

Slide 14: Legacy Path as Comparison

The old guidance-only path is still useful as a comparison benchmark and for reproducibility. But it should not be presented as the main contribution anymore, because the expectation measurement is weaker and the interpretation is less clean.

Slide 15: Final Conclusion

Our final conclusion is narrower and more honest. A Tushare-first earnings-surprise framework is more credible than the previous weak-proxy design. But the main contribution of the repository right now is a cleaner and auditable diagnostic baseline showing how expectation measurement and match quality affect inference.

Slide 16: Next Step

The next step is not to redesign the framework again. The next step is to keep the Tushare event architecture, preserve the focused diagnostic outputs, and only add stronger external expectation data if we want a stronger final headline.